

CS 2302 - Data Structures
Fall 2018
Project 3 - Option A

Overview

[Natural Language Processing \(NLP\)](#) is the sub-field of artificial intelligence that deals with designing algorithms, programs, and systems that can understand human languages in written and spoken forms.

[Word embeddings](#) are a recent advance in NLP that consists of representing words by vectors in such a way that if two words have similar meanings, their embeddings are also similar. See <https://nlp.stanford.edu/projects/glove/> for an overview of this interesting research.

In order to work in real-time, NLP systems such as Siri and Alexa need to efficiently retrieve the embeddings given their corresponding words. In this lab, you will implement a simple version of this. The web page mentioned above contains links to files that contain word embeddings of various lengths for various vocabulary sizes. Use the file "glove.6B.50d.txt" which contains word embeddings of length 50 for a very large number of words. Each line in the file starts with the word being described, followed by 50 floating point numbers that represent the word's vector description (the embedding). The words are ordered by frequency of usage, so "the" is the first word.

Your task for this lab is to write a program that does the following:

1. Read the file *glove.6B.50d.txt* and store each word and its embedding in a binary search tree. Ask the user what type of binary search tree he/she wants to use (AVL Tree or Red-Black Tree). You are free to use the implementation provided in your zyBook for these two types of trees. Adapt zyBook's code to include the word and its embedding and use the word as key. Ignore the "words" in the file that do not start with an alphabetic character (for example ",", and ".").
2. Read another file containing pairs of words (two words per line) and for every pair of words find and display the "similarity" of the words (see example in appendix). To find the similarity of words w_0 and w_1 , with embeddings e_0 and e_1 , we use the cosine distance, which ranges from -1 to 1, given by:

$$\text{sim}(w_0, w_1) = \frac{e_0 \cdot e_1}{|e_0| |e_1|}$$

Where $e_0 \cdot e_1$ is the dot product of e_0 and e_1 and $|e_0|$ and $|e_1|$ are the magnitudes of e_0 and e_1 . Look-up these formulas online if necessary.

3. Write the following methods to extract information from the tree.

- (a) Compute the number of nodes in the tree.
- (b) Compute the height of the tree.
- (c) Generate a file containing all the words stored in the tree, in ascending order, one per line.
- (d) Given a desired depth, generate a file with all the keys that have that depth, in ascending order.

Recall that the root has depth zero, its children have depth one, and so on.
As usual, write a report describing your work. Determine the $O()$ running times of your methods and show tables illustrating their actual running times and discuss disagreements between theoretical and experimental results.

Appendix

Word similarities:

barley shrimp 0.5352693296408768
barley oat 0.6695943991361065
federer baseball 0.28697851474308855
federer tennis 0.7167608209373065
harvard stanford 0.8466463229405593
harvard utep 0.06842559024883524
harvard ant -0.026703792920826475
raven crow 0.615012250504612
raven whale 0.32908915459219595
spain france 0.7909148906835685
spain mexico 0.7513763544646808
mexico france 0.5477963415949284

What you need to do

Part 1 - Due Friday, November 2, 2018

Implement the program described above, and upload your code to GitHub.

Part 2 - Due Tuesday, November 6, 2018

Add your team members as collaborators to your GitHub repo. They will add you to their projects as a collaborator as well. Read their code and give them feedback. Use *pull requests* and/or the *Issues* section to do so .

Extra Credit

Re-do the lab using a B-Tree, run multiple experiments using different values for the degree of the tree, and include in your report how performance changes as the degree varies. Finally, use tables and graphs to compare B-Trees with AVL and Red-Black trees.

Rubric

Criteria	Proficient	Neutral	Unsatisfactory
Correctness	The code compiles, runs, and solves the problem.	The code compiles, runs, but does not solve the problem (partial implementation).	The code does not compile/run, or little progress was made.
Space and Time complexity	Appropriate for the problem.	Can be greatly improved.	Space and time complexity not analyzed
Problem Decomposition	Operations are broken down into loosely coupled, highly cohesive methods	Operations are broken down into methods, but they are not loosely coupled/highly cohesive	Most of the logic is inside a couple of big methods
Style	Variables and methods have meaningful/appropriate names	Only a subset of the variables and methods have meaningful/appropriate names	Few or none of the variables and methods have meaningful/appropriate names
Robustness	Program handles erroneous or unexpected input gracefully	Program handles some erroneous or unexpected input gracefully	Program does not handle erroneous or unexpected input gracefully
Documentation	Non-obvious code segments are well documented	Some non-obvious code segments are documented	Few or none non-obvious segments are documented
Code Review	Useful feedback was provided to team members.	Feedback was provided to team members, but it was not very useful.	Little to no feedback was provided to team mates.

	Feedback received from team members was used to improve the code.	Feedback received from team mates was partially used to improve the code	Received feedback was not used to improve the code.
Report	Covers all required material in a concise and clear way with proper grammar and spelling.	Covers a subset of the required material in a concise and clear way with proper grammar and spelling.	Does not cover enough material and/or the material is not presented in a concise and clear way with proper grammar and spelling.